


Assignment Project Exam Help

Add WeChat powcoder



Supervised Learning I: Regression

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Today

Add WeChat powcoder

- Multivariate linear regression
- Solution for SSD cost
 - Indirect
 - Direct
- Maximum likelihood cost

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help


Linear Regression

Add WeChat powcoder

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

θ_i 's: Parameters <https://powcoder.com>



Cost Function: Add WeChat powcoder

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

SSD = sum of squared differences, also

SSE = sum of squared errors

Assignment Project Exam Help

Multidimensional inputs

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

Notation:

n = number of features

$x^{(i)}$ = input (features) of i^{th} training example.

$x_j^{(i)}$ = value of feature j in i^{th} training example.

Assignment Project Exam Help

Multivariate Linear Regression

Add WeChat powcoder

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

Assignment Project Exam Help

For convenience of notation, define $x_0 = 1$.

θ_i 's: Parameters <https://powcoder.com>

Cost Function: Add WeChat powcoder

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(\theta_0, \theta_1, \dots, \theta_n)$ **How??**
 $\theta_0, \theta_1, \dots, \theta_n$

Assignment Project Exam Help

Two potential solutions

$$\min_{\theta} J(\theta; x^{(1)}, y^{(1)}, \dots, x^{(m)}, y^{(m)})$$

Gradient descent (iterative algorithm)


- Start with a guess for θ
- Change θ to decrease $J(\theta)$
- Until reach minimum

Direct minimization

- Take derivative, set to zero
- Sufficient condition for minima
- Not possible for most “interesting” cost functions

Assignment Project Exam Help

Add WeChat powcoder



Solving Linear
Assignment Project Exam Help
Regression

<https://powcoder.com>

Add WeChat powcoder

Gradient Descent

Assignment Project Exam Help

Gradient Descent Algorithm

Set $\theta = 0$

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad \text{simultaneously for all } j = 0, \dots, n$$


} until convergence

Assignment Project Exam Help Gradient Descent: Intuition

Add WeChat powcoder


$$h_{\theta}(x)$$

(for fixed θ_1 , this is a function of x)



$$J(\theta_1)$$

(function of the parameter θ_1)




$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Assignment Project Exam Help Gradient Descent: Intuition

Add WeChat powcoder


$$h_{\theta}(x)$$

(for fixed θ_1 , this is a function of x)




$$J(\theta_1)$$

(function of the parameter θ_1)



$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Assignment Project Exam Help



Assignment Project Exam Help

Gradient for Least Squares Cost

Add WeChat powcoder

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Assignment Project Exam Help

$$\frac{\partial}{\partial \theta_j} J(\theta) =$$

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Gradient for Least Squares Cost

Add WeChat powcoder

For one example

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Gradient for Least Squares Cost

Add WeChat powcoder

For one example

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2 \\ &= 2 \cdot \frac{1}{2} (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y)\end{aligned}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Gradient for Least Squares Cost

Add WeChat powcoder

For one example

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2 \\&= 2 \cdot \frac{1}{2} (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y) \\&= (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} \left(\sum_{i=0}^n \theta_i x_i - y \right)\end{aligned}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Gradient for Least Squares Cost

Add WeChat powcoder

For one example

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2 \\&= 2 \cdot \frac{1}{2} (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y) \\&= (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} \left(\sum_{i=0}^n \theta_i x_i - y \right) \\&= (h_{\theta}(x) - y) x_j\end{aligned}$$

What is this?

Assignment Project Exam Help

Gradient Descent Algorithm

Set $\theta = 0$

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) \mathbf{x}_j^{(i)}$$

simultaneously for all $j = 0, \dots, n$

} until convergence

Add WeChat powcoder

in vector form?

$$\theta := \theta + \alpha \frac{1}{m} \mathbf{e}^T \mathbf{X}$$

$$\mathbf{e}_i = h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}$$

$$\mathbf{X} = \begin{bmatrix} x_1^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$$

Assignment Project Exam Help

Feature normalization


Add WeChat powcoder

- If features have very different scale, GD can get “stuck” since x_j affects size of gradient in the direction of j^{th} dimension
- Normalizing features to be zero-mean (μ) and same-variance (σ) helps gradient descent converge faster

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder




$J(\theta)$ normalize

$$x_j^{(i)} = \frac{x_j^{(i)} - \mu_j}{\sigma_j}$$



Assignment Project Exam Help

Add WeChat powcoder



Assignment Project Exam Help

Solving Linear Regression

<https://powcoder.com>

Add WeChat powcoder

Direct Solution

Assignment Project Exam Help

Direct solution

Add WeChat powcoder

Want to minimize SSD:

$$J(\theta_0, \theta_1, \dots, \theta_m) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Assignment Project Exam Help


Find minima of function <https://powcoder.com>

$$\theta \in \mathbb{R}^{n+1}$$

Add WeChat powcoder

$$\frac{\partial}{\partial \theta_j} J(\theta) = \dots = 0 \quad (\text{for every } j)$$

Solve for $\theta_0, \theta_1, \dots, \theta_n$



Assignment Project Exam Help

Direct solution

Add WeChat powcoder

Re-write SSD using vector-matrix notation:

$$J(\theta) = \frac{1}{2m} (X\theta - y)^T (X\theta - y)$$


Assignment Project Exam Help

Where:

<https://powcoder.com>

$$X = \begin{bmatrix} - & (x^{(1)})^T & - \\ - & (x^{(2)})^T & - \\ & \vdots & \\ - & (x^{(m)})^T & - \end{bmatrix} \quad \bar{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

Add WeChat powcoder



Solution: Normal Equation

$$\theta = (X^T X)^{-1} X^T y$$

Assignment Project Exam Help

Derivation of Normal Equations

Add WeChat powcoder

- SSE in matrix form:

$$\begin{aligned} J(\theta) &= \frac{1}{2m} (X\theta - y)^T (X\theta - y) = \\ &= \frac{1}{2m} \{ \theta^T \{X^T X\} \theta - 2\{X^T y\}^T \theta + \text{const} \} \end{aligned}$$

<https://powcoder.com>

- Take derivative with respect to θ (vector), set to 0

$$\frac{\partial J}{\partial \theta} \propto X^T X \theta - X^T y = 0$$

ignore constant multiplier

$$\theta = (X^T X)^{-1} X^T y$$

- Also known as the **least mean squares**, or **least squares** solution

Assignment Project Exam Help

Example: $m = 4$.

Add WeChat powcoder

	Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
x_0	x_1	x_2	x_3	x_4	y
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Design
Matrix

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}$$

$$y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

Normal
Equation

$$\theta = (X^T X)^{-1} X^T y$$

Assignment Project Exam Help

Trade-offs

Add WeChat powcoder

m training examples, n features.

Gradient Descent

Normal Equations

Assignment Project Exam Help


- | | |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none">• Need to choose α.• Needs many iterations.• Works well even when n is large. | <ul style="list-style-type: none">• No need to choose α.• Don't need to iterate.• Need to compute $(X^T X)^{-1}$• Slow if n is very large. |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Add WeChat powcoder



Maximum Likelihood for Linear Regression

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help
So far, we have treated outputs as
noiseless
Add WeChat powcoder

- Defined cost function as “distance to true output” Assignment Project Exam Help
- An alternate view:
<https://powcoder.com>
Add WeChat powcoder
 - data (x, y) are generated by unknown process
 - however, we only observe a noisy version
 - how can we model this uncertainty?
- Alternative cost function?

Assignment Project Exam Help

How to model uncertainty in data?

Add WeChat powcoder

Hypothesis:

$$h_{\theta}(x) = \theta^T x$$

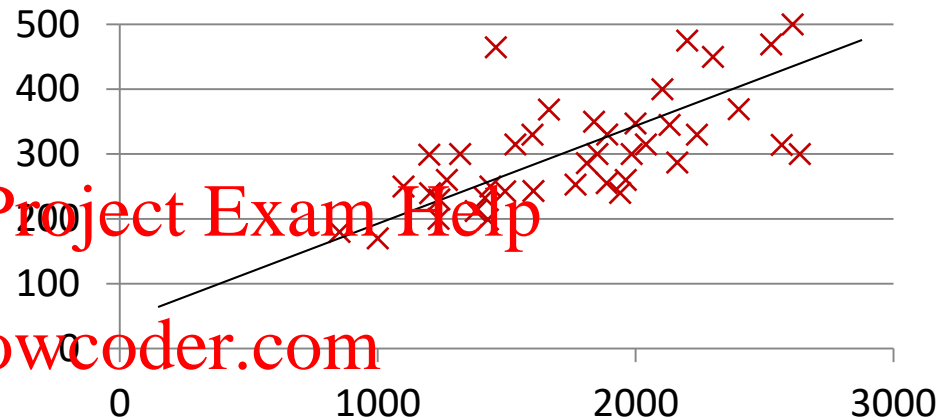
θ : parameters

$D = (x^{(i)}, y^{(i)})$: data

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



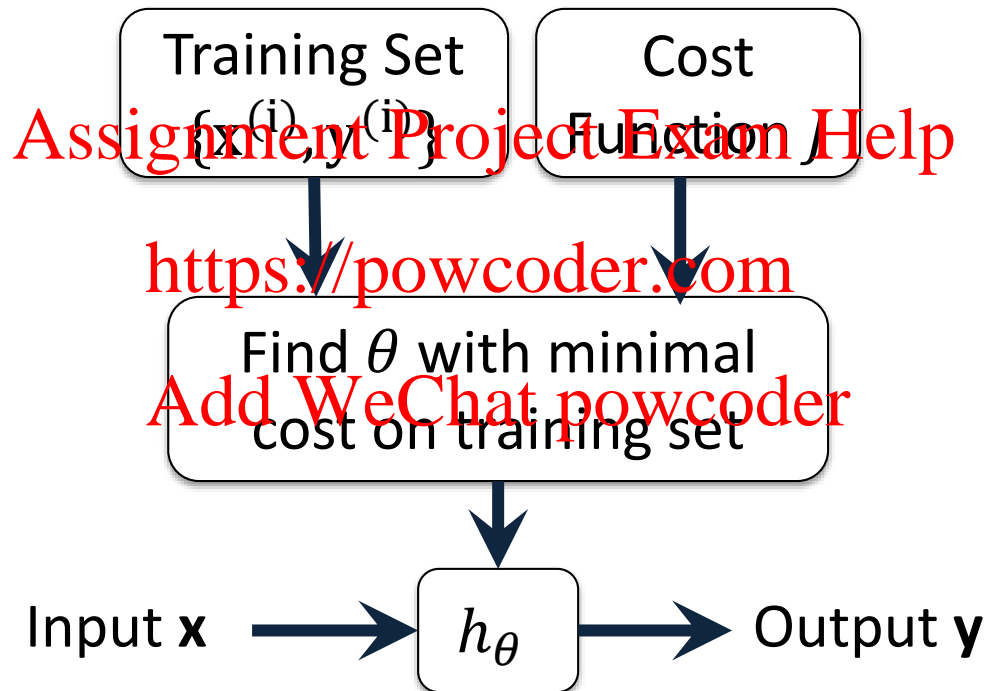
New cost function:

maximize probability of data given model:

$$p((x^{(i)}, y^{(i)}) | \theta)$$


Recall: Cost Function

Add WeChat powcoder



Assignment Project Exam Help

Alternative View: "Maximum Likelihood"



Assignment Project Exam Help

Maximum Likelihood: Example

- Intuitive example: Estimate a coin toss

I have seen 3 flips of heads, 2 flips of tails, what is the chance of head (or tail) on my next flip?

- Model: <https://powcoder.com>

Each flip is a Bernoulli random variable X

X can take only two values: 1 (head), 0 (tail)

$$p(X = 1) = \theta, \quad p(X = 0) = 1 - \theta$$

- θ is a parameter to be identified from data

Maximum Likelihood: Example

Add WeChat powcoder

- 5 (independent) trials



$$X_1 = 1$$



$$X_2 = 0$$



$$X_3 = 1$$



$$X_4 = 1$$



$$X_5 = 0$$

Assignment Project Exam Help

- Likelihood of all 5 observations:

<https://powcoder.com>

Add WeChat powcoder

$$p(X_1, \dots, X_5 | \theta) = \theta^3 (1 - \theta)^2$$

- Intuition

ML chooses θ such that likelihood is maximized

Maximum Likelihood: Example

Add WeChat powcoder

- 5 (independent) trials



$$X_1 = 1$$



$$X_2 = 0$$



$$X_3 = 1$$



$$X_4 = 1$$



$$X_5 = 0$$

Assignment Project Exam Help

- Likelihood of all 5 observations:

<https://powcoder.com>

Add WeChat powcoder

$$p(X_1, \dots, X_5 | \theta) = \theta^3 (1 - \theta)^2$$

- Solution (left as exercise)

$$\theta_{ML} = \frac{3}{(3 + 2)}$$

i.e. fraction of heads in total number of trials

Assignment Project Exam Help

PSet 1 Out

Add WeChat powcoder

- Due on Tuesday 9/15 11:59pm GMT -5
(Boston Time)

Assignment Project Exam Help

- Diagnostic homework covering topics covered
in prereqs <https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help

Next Class

Add WeChat powcoder

Supervised Learning II: Classification:

classification; sigmoid function; logistic regression.

Assignment Project Exam Help

<https://powcoder.com>

Reading: Bishop 4.3.1-4.3.2; 4.3.4

Add WeChat powcoder

[overview of logistic regression](#)