

Assignment Project Exam Help

# Announcements

Add WeChat powcoder

**Reminder:** Class challenge out! Ends December 10th

Assignment Project Exam Help

- Final in two weeks, practice question will be posted today

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

## Class Challenge Task 2

Add WeChat powcoder

- 100 labels isn't enough! (see Lecture 21/first half of 22)

Assignment Project Exam Help

- Transfer learning from a different dataset? if it isn't *a dataset we provided*, you have to check with us (ImageNet OK)

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Ethics in Machine Learning

Kate Saenko

CS 542 Machine Learning



# Assignment Project Exam Help

AI Fears:

Which of these has already happened?



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

A robot kills human

yes, Uber car accident

AI takes over our lives

yes? Youtube algorithm

AI is watching us

yes, virtual police lineup

Autonomous weapons

not yet?

Humans losing jobs

Yes, e.g. librarians



# Assignment Project Exam Help

## AI Fears

Add WeChat powcoder

- **Autonomous weapons** – frameworks for regulation
- **Future of work** – deskilling / reskilling / superskilling.  
Many jobs that currently seem least likely to be automated have been racialized and gendered in ways connected to care and immigration, and have rarely paid living wages; how do we revalue work?
- **Worse Inequality** – digital divides; bias in algorithms may worsen inequality; ecological concerns in energy, storage and cooling required for ML; economic inequality
- **Divided societies with algorithmic bubbles** – challenges of populism, automated recommendations, news feeds; deepfakes and election meddling

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





Assignment Project Exam Help

Add WeChat powcoder

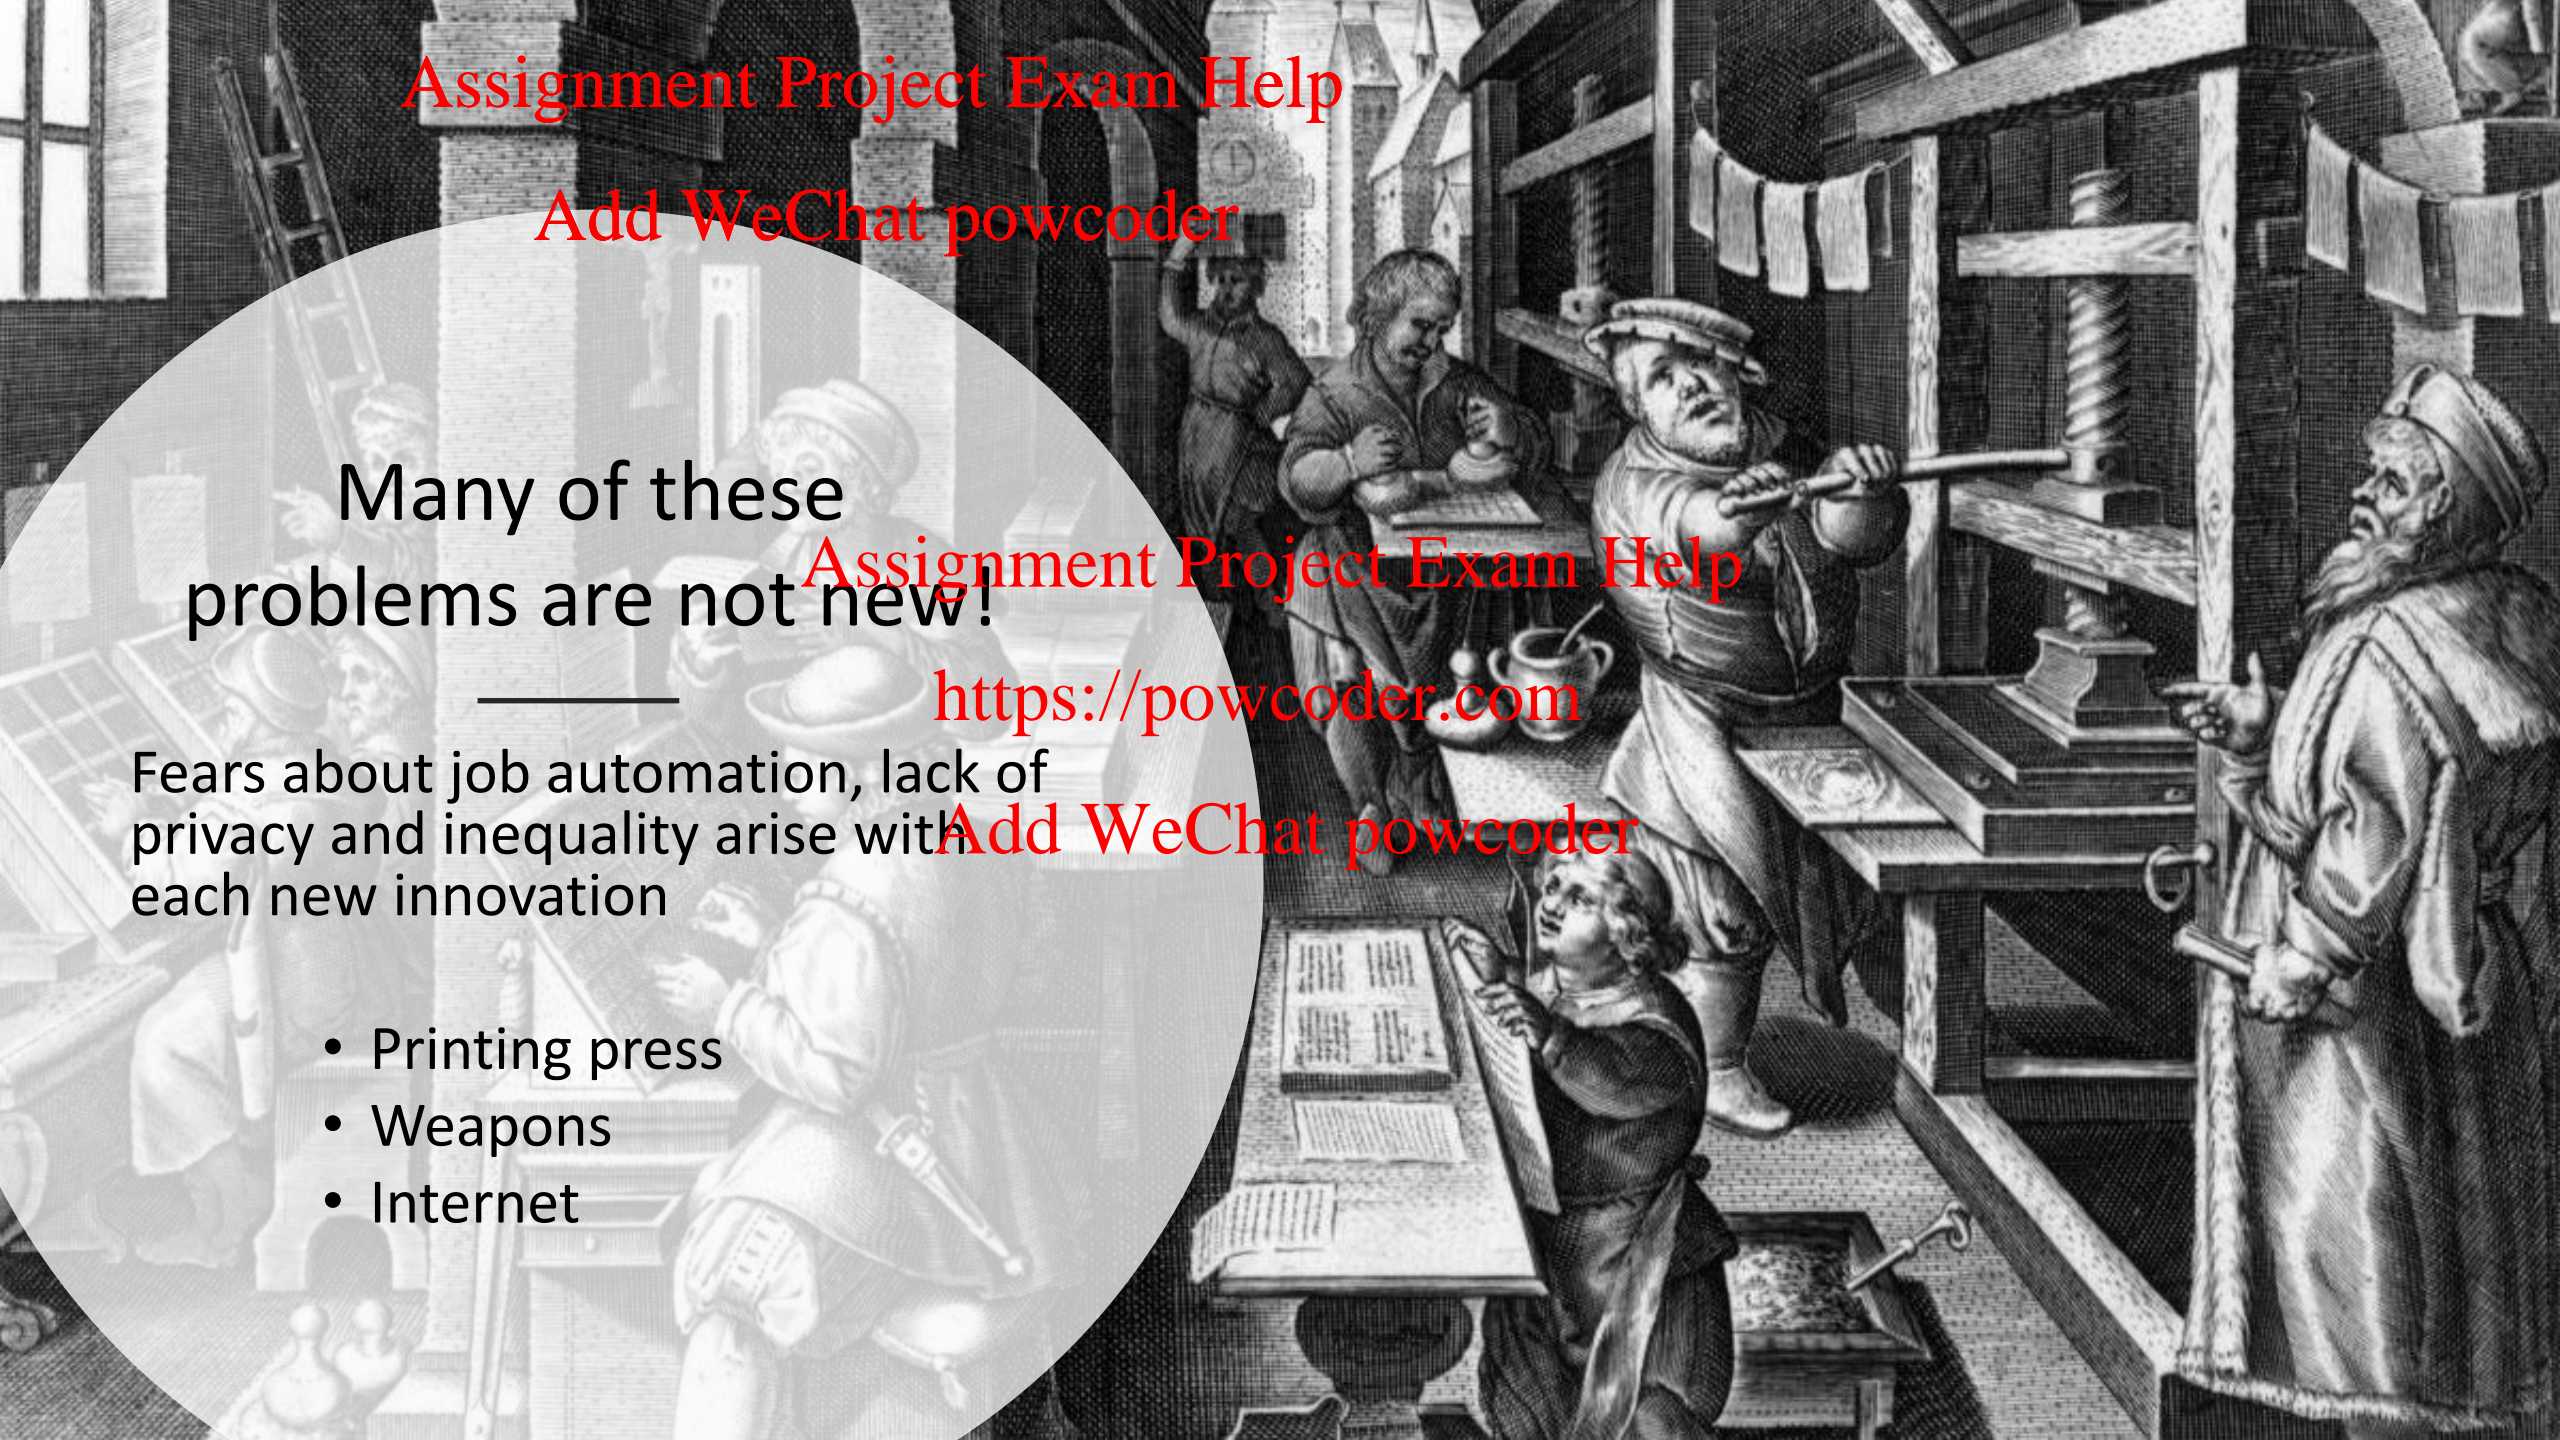
Many of these  
problems are not new!

Assignment Project Exam Help

<https://powcoder.com>

Fears about job automation, lack of  
privacy and inequality arise with  
each new innovation

- Printing press
- Weapons
- Internet





# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance

Assignment Project Exam Help

Add WeChat powcoder

Report: [A.I. would eliminate 75 million jobs but may create about 130 million jobs globally.](#)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Stable Roles	New Roles	Redundant Roles
Managing Directors and Chief Executives	Data Analysts and Scientists*	Data Entry Clerks
General and Operations Managers*	AI and Machine Learning Specialists	Accounting, Bookkeeping and Payroll Clerks
Software and Applications Developers and Analysts*	General and Operations Managers*	Administrative and Executive Secretaries
Data Analysts and Scientists*	Big Data Specialists	Assembly and Factory Workers
Sales and Marketing Professionals*	Digital Transformation Specialists	Client Information and Customer Service Workers*
Sales Representatives, Wholesale and Manufacturing, Technical and Scientific Products	Sales and Marketing Professionals*	Business Services and Administration Managers
Human Resources Specialists	New Technology Specialists	Accountants and Auditors
Financial and Investment Advisers	Organizational Development Specialists*	Material-Recording and Stock-Keeping Clerks
Database and Network Professionals	Software and Applications Developers and Analysts*	General and Operations Managers*
Supply Chain and Logistics Specialists	Information Technology Services	Postal Service Clerks
Risk Management Specialists	Process Automation Specialists	Financial Analysts
Information Security Analysts*	Innovation Professionals	Cashiers and Ticket Clerks
Management and Organization Analysts	Information Security Analysts*	Mechanics and Machinery Repairers
Electrotechnology Engineers	Ecommerce and Social Media Specialists	Telemarketers
Organizational Development Specialists*	User Experience and Human-Machine Interaction Designers	Electronics and Telecommunications Installers and Repairers
Chemical Processing Plant Operators	Training and Development Specialists	Bank Tellers and Related Clerks
University and Higher Education Teachers	Robotics Specialists and Engineers	Car, Van and Motorcycle Drivers
Compliance Officers	People and Culture Specialists	Sales and Purchasing Agents and Brokers
Energy and Petroleum Engineers	Client Information and Customer Service Workers*	Door-To-Door Sales Workers, News and Street Vendors, and Related Workers
Robotics Specialists and Engineers	Service and Solutions Designers	Statistical, Finance and Insurance Clerks
Petroleum and Natural Gas Refining Plant	Digital Marketing and Strategy Specialists	Lawyers



# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance

Assignment Project Exam Help

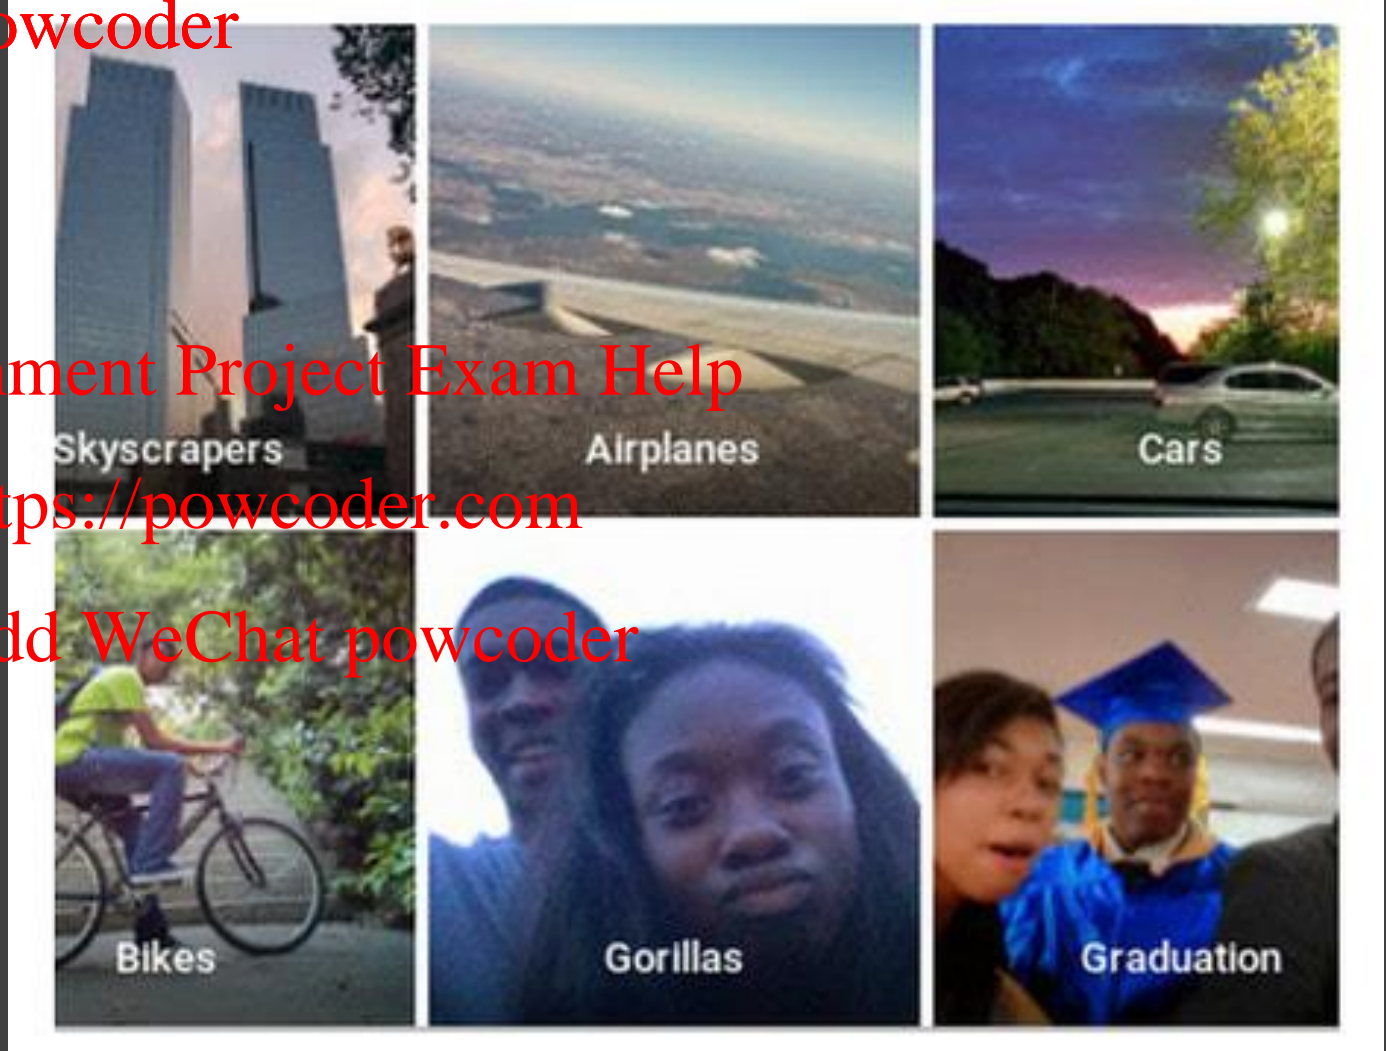
Bias can lead to offensive or unfair results...

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




# Assignment Project Exam Help

## Gender Shades (Buolamwini & Gebru, 2018)

### Add WeChat powcoder

- Evaluated commercial gender classifiers from Microsoft, FACE++, IBM
- Found large disparity in error between population subgroups based on gender, skin color

### Assignment Project Exam Help

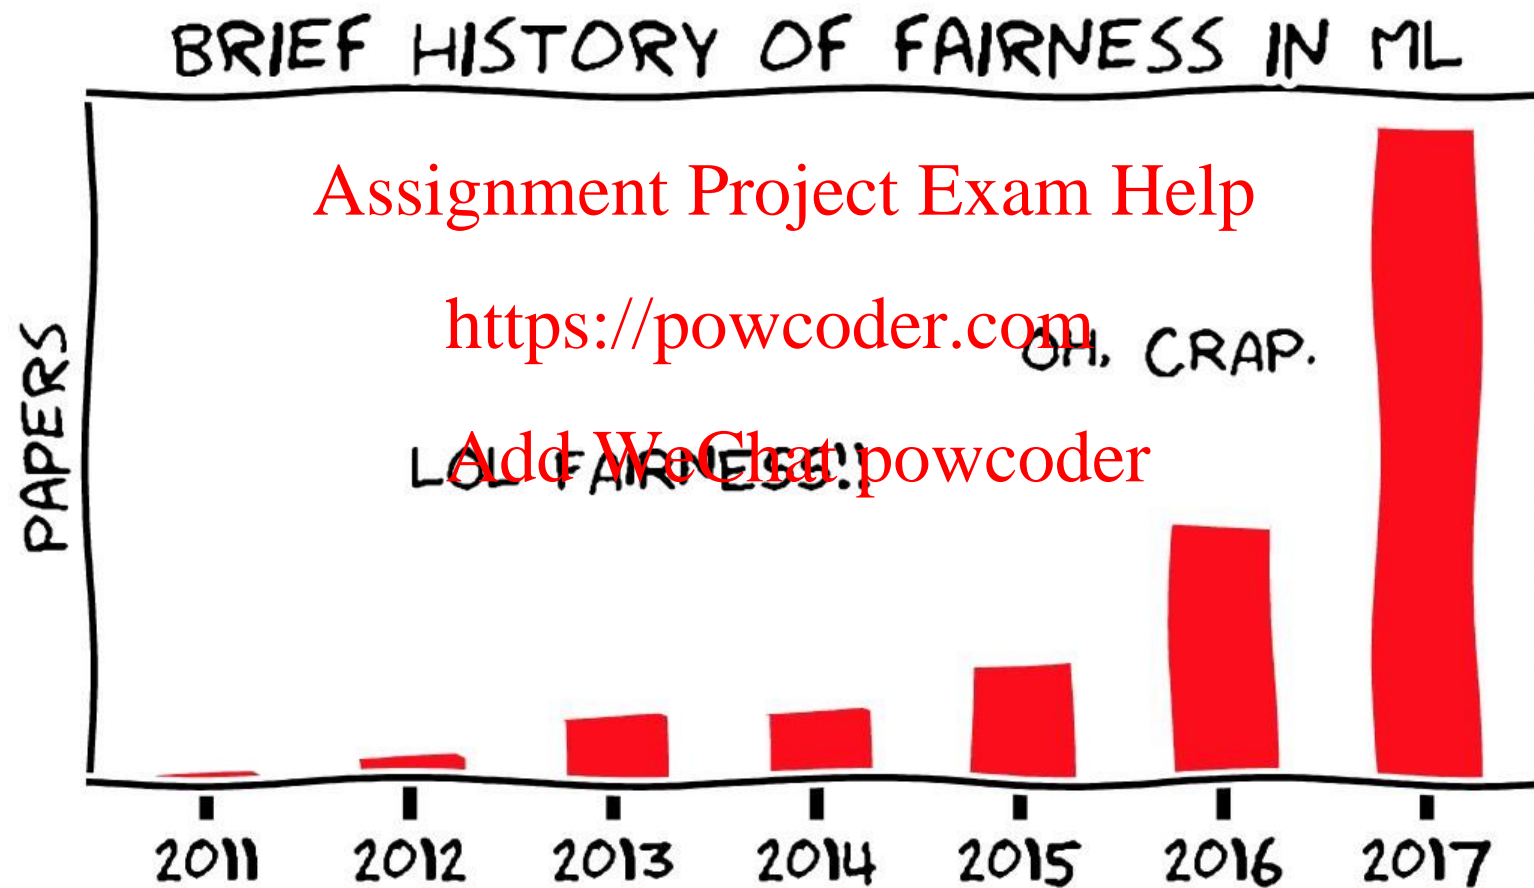


Buolamwini, Joy, and Timnit Gebru. "Gender shades: Intersectional accuracy disparities in commercial gender classification." Conference on Fairness, Accountability and Transparency. 2018.



# Assignment Project Exam Help

Add WeChat powcoder Fairness in Machine Learning



Moritz Hardt

# Assignment Project Exam Help Initially: AI is better than humans!

Add WeChat powcoder

Can an Algorithm Hire Better Than a Human?



Claire Cain Miller @clairecm

The Algorithm That Beats Your Bank Manager

PREDICTIVE POLICING: USING  
MACHINE LEARNING  
PATTERNS OF CRIME

Assignment Project Exam Help



The Marshall Project Nonprofit journalism about criminal justice

SEARCH

ABOUT

DONATE



<https://powcoder.com>

## The New Science of Sentencing

*Should prison sentences be based on new data? Haven't been committed yet?*

Add WeChat powcoder





Assignment Project Exam Help  
Wait, maybe not such a good idea...

Add WeChat powcoder

## Beauty contest judged by AI and the robots discriminate against dark skin

3 days ago | Published by : Avinash Nandakumar

Is an algorithm any less racist than a human?



CNN Money U.S. + Business Markets Tech Media Personal Finance Small Biz Luxury stock tickers

## Math is racist: How data is driving inequality

by Aimee Rawlins @aimeerawlins

Assignment Project Exam Help

## Machine Bias

<https://powcoder.com>

Add WeChat powcoder

HIDDEN BIAS

## When Algorithms Discriminate



Claire Cain Miller @clairecm JULY 9, 2015

The online world is shaped by forces beyond our control, determining the stories we read on Facebook, the people we meet on OkCupid and the search results we see on Google. Big data is used to make decisions about health care, employment, housing, education and policing.

But can computer programs be discriminatory?

There is a widespread belief that software and algorithms that rely on data are objective. But software is not free of human influence. Algorithms are written and maintained by people, and machine learning algorithms adjust what they do based on people's behavior. As a result, say researchers in computer science, ethics and law, algorithms can reinforce human prejudices.

Google's online advertising system, for instance, showed an ad for high-income jobs to men much more often than it showed the ad to women, a new study by Carnegie Mellon University researchers found.



Photo by Ben Torres (Bloomberg)

## Big Bad Data May Be Triggering Discrimination

August 15, 2016

AUTHORS

**B** Bloomberg BNA - Staff Reports

SHARING

Twitter

By Kevin McGowan, Bloomberg BNA

"Big data" is filled with promise for improving recruitment and hiring, but if employers don't take care it can also drive them to unintentionally commit discrimination.

"It's a bit of a black box," said Commissioner Victoria Lipnic (R) of the Equal Employment Opportunity Commission, referring to the formulas data analysts and programmers develop

# Example of ML (un)fairness: COMPAS

Add WeChat powcoder

- Criminal justice: recidivism algorithms (COMPAS)
- Predicting if a defendant should receive bail
- Unbalanced false positive rates: more likely to wrongly deny a black person bail

Assignment Project Exam Help

<https://powcoder.com>

ProPublica Analysis of COMPAS Algorithm

Add WeChat powcoder

	White	Black
Wrongly Labeled High-Risk	23.5%	<b>44.9%</b>
Wrongly Labeled Low-Risk	<b>47.7%</b>	28.0%



# Assignment Project Exam Help

Add WeChat powcoder

## Two Petty Theft Arrests

**VERNON PRATER**

### Prior Offenses

2 armed robberies, 1  
attempted armed  
robbery

### Subsequent Offenses

1 grand theft

**LOW RISK**

**3**

**BRISHA BORDEN**

### Prior Offenses

4 juvenile  
misdemeanors

### Subsequent Offenses

None

**HIGH RISK**

**8**

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Example of ML (un)fairness: word embedding


Add WeChat powcoder

- Bias found in word embeddings (Bolukbasi et al. 2016)
  - Examined word embeddings (word2vec) trained on Google News
  - Represent each word with high-dimensional vector
  - Vector arithmetic: found analogies like
  - Paris - France = London - England
  - man - woman = programmer - homemaker = surgeon - nurse
- The good news: word embeddings learn so well!
- The bad news: sometimes too well
- Our chatbots should be less biased than we are

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





## Assignment Project Exam Help

# Example of ML (un)fairness: word embedding

Extreme <i>she</i>	Extreme <i>he</i>	Gender stereotype <i>she-he</i> analogies		
1. homemaker	1. maestro	sewing-carpentry	registered nurse-physician	housewife-shopkeeper
2. nurse	2. skipper	nurse-surgeon	interior designer-architect	softball-baseball
3. receptionist	3. protege	blond-burly	feminism-conservatism	cosmetics-pharmaceuticals
4. librarian	4. philosopher	giggle-chuckle	vocalist-guitarist	petite-lanky
5. socialite	5. captain	sassy-snappy	diva-superstar	charming-affable
6. hairdresser	6. architect	volleyball-football	cupcakes-pizzas	lovely-brilliant
7. nanny	7. financier	Gender appropriate <i>she-he</i> analogies		
8. bookkeeper	8. warrior	queen-king	sister-brother	mother-father
9. stylist	9. broadcaster	waitress-waiter	ovarian cancer-prostate cancer	convent-monastery
10. housekeeper	10. magician			

TABLE. Left: The most extreme occupations as projected on to the she he gender direction on w2vNEWS. Occupations such as businesswoman, where gender is suggested by the orthography, were excluded. Right: Automatically generated analogies for the pair she-he using the procedure in paper.

## Assignment Project Exam Help

# Machine Learning and Social Norms

- Sample norms: privacy, fairness, transparency, accountability...
- Possible approaches
  - “traditional”: legal, regulatory, watchdog
  - *Embed* social norms in data, algorithms, models
- Case study: privacy-preserving machine learning
  - “single”, strong, definition (differential privacy)
  - almost every ML algorithm has a private version
- Fair machine learning
  - not so much...
  - impossibility results

# Assignment Project Exam Help

## (Un)Fairness Where?

Add WeChat powcoder

- Data (input)
  - e.g. more arrests where there are more police
  - Label should be “committed a crime”, but is “convicted of a crime”
  - try to “correct” bias
- Models (output)
  - e.g. discriminatory treatment of subpopulations
  - build or “post-process” models with subpopulation guarantees
  - equality of false positive/negative rates; calibration
- Algorithms (process)
  - learning algorithm *generating* data through its decisions
  - e.g. don’t learn outcomes of denied mortgages
  - lack of clear train/test division
  - design (sequential) *algorithms* that are fair

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Data Bias Examples

Add WeChat powcoder

## Reporting Bias example:

A sentiment-analysis model is trained to predict whether book reviews are positive or negative based on a corpus of user submissions to a popular website. The majority of reviews in the training data set reflect extreme opinions (reviewers who either loved or hated a book), because people were less likely to submit a review of a book if they did not respond to it strongly. As a result, the model is less able to correctly predict sentiment of reviews that use more subtle language to describe a book.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## Selection Bias example:

A model is trained to predict future sales of a new product based on phone surveys conducted with a sample of consumers who bought the product. Consumers who instead opted to buy a competing product were not surveyed, and as a result, this group of people was not represented in the training data.

# Why fairness is hard

Add WeChat powcoder

- Suppose we are a bank trying to fairly decide who should get a loan i.e. Who is most likely to pay us back?

Assignment Project Exam Help

- Suppose we have two groups, A and B (the sensitive attribute)  
This is where discrimination could occur

<https://powcoder.com>

Add WeChat powcoder

- The simplest approach is to remove the sensitive attribute from the data, so that our classifier doesn't know the sensitive attribute

## Assignment Project Exam Help

# Why fairness is hard

Add WeChat powcoder

Table 2: To Loan or Not to Loan?

Age	Gender	Postal Code	Req Amt	A or B?	Pay
46	F	M5E	\$300	A	1
24	M	M4C	\$1000	B	1
33	M	M3H	\$250	A	1
34	F	M9C	\$2000	A	0
71	F	M3B	\$200	A	0
28	M	M5W	\$1500	B	0



# Why fairness is hard

Add WeChat powcoder

Table 3: To Loan or Not to Loan? (masked)

Age	Gender	Postal Code	Req Amt	A or B?	Pay
46	F	M5E	\$300	?	1
24	M	M4C	\$1000	?	1
33	M	M3H	\$250	?	1
34	F	M9C	\$2000	?	0
71	F	M3B	\$200	?	0
28	M	M5W	\$1500	?	0

# Why fairness is hard

Add WeChat powcoder

- However, if the sensitive attribute is correlated with the other attributes, this isn't good enough

Assignment Project Exam Help

<https://powcoder.com>

- It is easy to predict race if you have lots of other information (e.g. home address, spending patterns)

Add WeChat powcoder

- More advanced approaches are necessary


# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

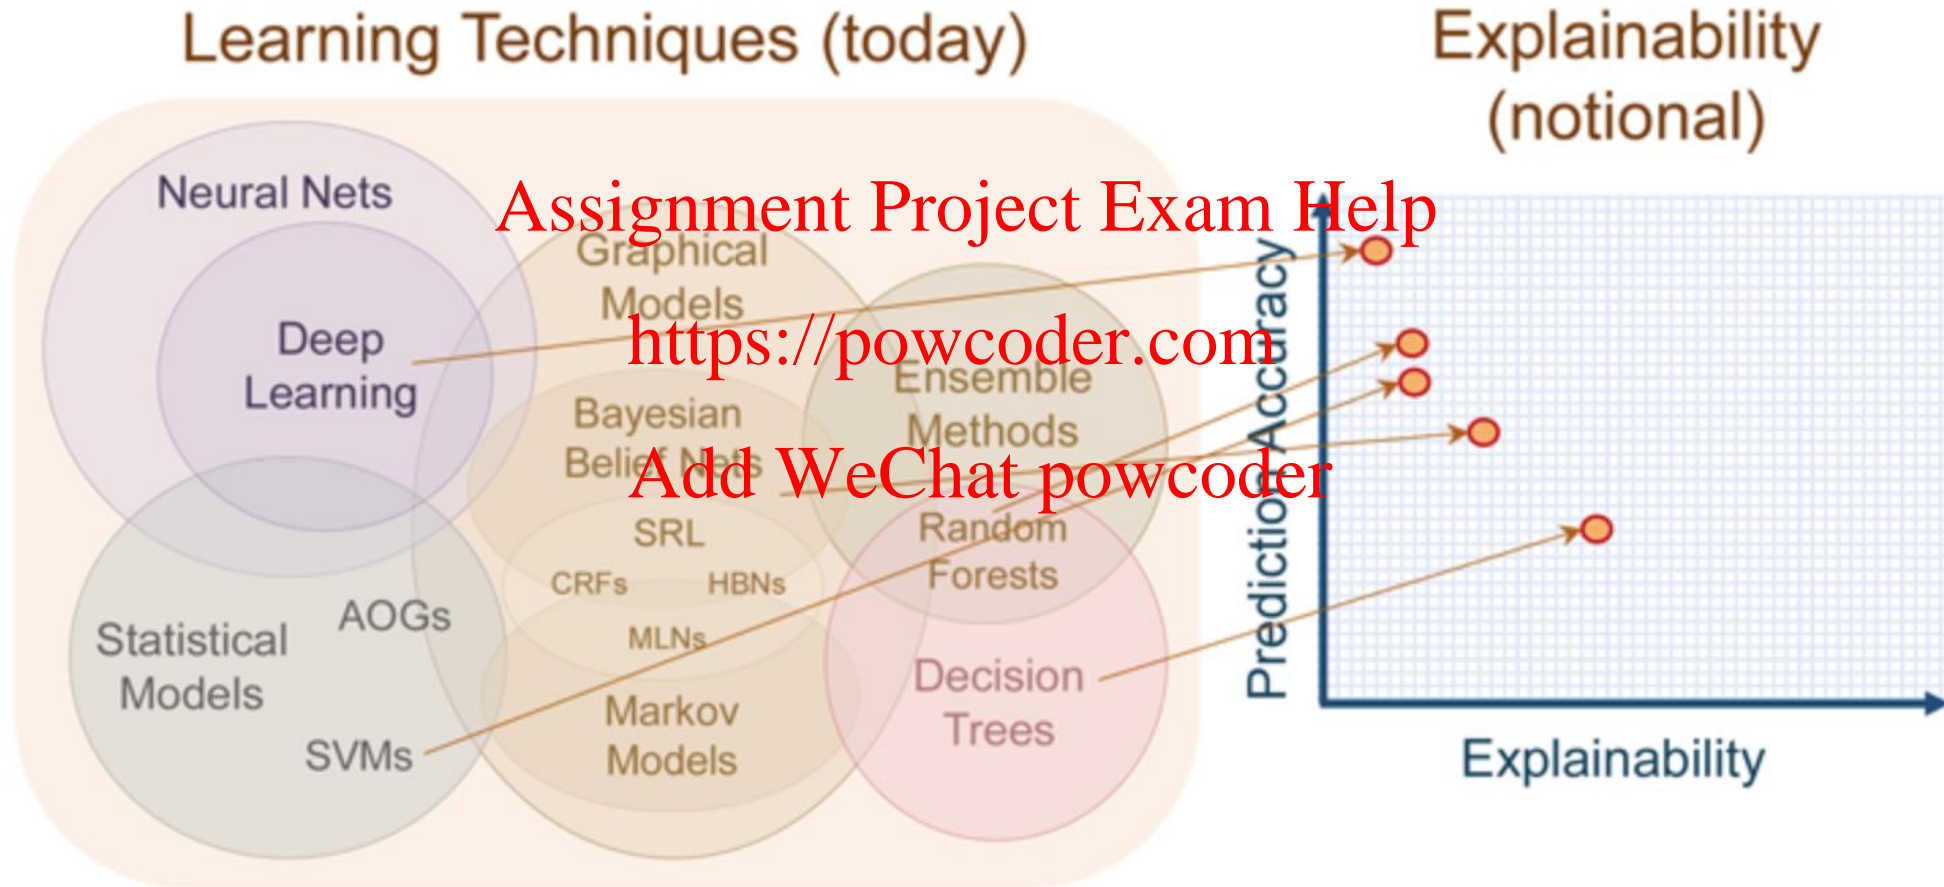




Assignment Project Exam Help

# Accuracy vs explainability

Add WeChat powcoder



E.g.: dataset bias leads to higher errors on 'novel' data...  
Can an explanation point to such bias?

## Training

Most cows are black/brown



Most sheep are white



## Test

Prediction: "cow" 76%



True class: "sheep"

## Explanation



# Assignment Project Exam Help


## Gender bias in captioning models (Hendricks et al. 2018)

**Evidence for "man"**




Baseline: A **man** sitting at a desk with a laptop computer.

### Ground truth captions



### Generated captions





# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- **AI Supremacy**
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance

Assignment Project Exam Help

If we start trusting algorithms to make decisions, who will have the final word on important decisions? Will it be humans, or algorithms?

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Algorithms are already being used to determine prison sentences. Judges' decisions are affected by their moods, so some argue that judges should be replaced with "robojudges". However, a ProPublica study found that one of these popular sentencing algorithms was highly biased against blacks.

# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



<https://www.youtube.com/watch?v=VWrhRBb-1lg>

# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- **Autonomous weapons**
- Self-driving cars
- Privacy and surveillance



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



*The Centre for New American security said in a report that the Chinese company Ziyan is negotiating to sell Blowfish A2, a killer robot capable of 60 millimeter mortar shells or a 35-40 millimeter grenade launcher, to the governments of Pakistan and Saudi Arabia*





# Assignment Project Exam Help

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

**INCREASE IN VIOLENT CRIME**

**SDN**

# Assignment Project Exam Help

## Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- **Self-driving cars**
- Privacy and surveillance


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

In a preliminary report about the crash released in May, the National Transportation Safety Board said the Uber car's computer system had spotted Ms. Herzberg six seconds before impact, but classified Ms. Herzberg, who was not in a crosswalk, first as an unrecognized object, then as another vehicle and finally as a bicycle.

The death of Elaine Herzberg (August 2, 1968 – March 18, 2018) was the first recorded case of a pedestrian fatality involving a self-driving (autonomous) car, after a collision ... Following the fatal incident, Uber suspended testing of self-driving vehicles in Arizona, where such testing had been sanctioned since August 2016





# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance





# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance

One in two American adults is in a law enforcement face recognition network-- <https://www.perpetuallineup.org/>

Add WeChat powcoder  
Most law enforcement agencies do little to ensure their systems are accurate.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Assignment Project Exam Help


# Private data leakage in language models

Add WeChat powcoder

<https://talktotransformer.com/>

## Talk to Transformer

See how a modern neural network completes your text. Type a custom snippet or try one of the examples. [Learn more](#) below.

 Follow @AdamDanielKing for more neat neural networks.

At >\$2500 per month, I can no longer sustain this site without a sponsor. If you'd like to present your product or service here, please [email me](#).

Sci-fi

The spaceship was flying erratically, something was wrong...

GENERATE ANOTHER

## Generated Completion

The spaceship was flying erratically, something was wrong...

The whole ordeal left some thinking that the Guardians were taking orders from Slitheen or some alien facsimile of him. In truth, Guardians are commanded by the T-Shirt Guy, who has to have trouble controlling his fiery temper. Whether it's a direct order from Slitheen, just a telepathic suggestion, or some combination of both, the Guardians keep a watchful eye on the Other Man.

At the start of the episode, a mysterious Doctor tells us that a ship sent from Arcturus Minor

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

# Private data leakage in language models

Add WeChat powcoder

<https://talktotransformer.com/>

## Talk to Transformer

See how a modern neural network completes your text. Type a custom snippet or try one of the examples. [Learn more](#) below.



Follow @AdamDanielKing

for more neat neural networks.

At >\$2500 per month, I can no longer sustain this site without a sponsor. If you'd like to present your product or service here, please [email me](#).

Custom prompt

social security number (SSN):

GENERATE ANOTHER

## Generated Completion

social security number (SSN): XXX-XXX-XXXX

Department of Work and Pensions, Social Security Number program, Fax: 1-800-772-1213, email: chelseas@dualcredits.org or spousal number(s) listed in public records available on GOV.UK.

Social Security number

If you are applying online, you can enter your FFEA number instead of the SSN.

If you are applying from outside the UK, an application form is provided to you.

\*actual SSN was generated, we blocked it out



# Ethical Issues in Machine Learning

- Job Loss
- Algorithmic Bias
- Transparency
- AI Supremacy
- Fake news and videos
- Autonomous weapons
- Self-driving cars
- Privacy and surveillance

Further reading:

FATML Conference: <https://www.fatml.org/>

ACM FAT\* Conference: <https://fatconference.org/>

**Fairness, Accountability,  
and Transparency  
in Machine Learning**

ACM Conference on Fairness,  
Accountability, and Transparency  
(ACM FAT\*)

A computer science conference with a cross-disciplinary focus that brings together researchers and practitioners interested in fairness, accountability, and transparency in socio-technical systems.



Assignment Project Exam Help

Add WeChat powcoder

Next Class

## Applications III: On-device Machine Learning

knowledge distillation, pruning, quantization

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder